

Lab6

57118232 谢隆文

Task 1: Implementing a Simple Firewall

Task1.A Implement a Simple Kernel Module

原始目录存在空格，目录的空格被 make 识别为编译的 target，所以我们需要把 kernel_module 拷贝到 /home/seed/ 目录下进行编译。

```
[07/26/21]seed@VM:~/kernel_module$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  Building modules, stage 2.
  MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/kernel_module/hello.o
see include/linux/module.h for more information
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/26/21]seed@VM:~/kernel_module$ sudo insmod hello.ko
[07/26/21]seed@VM:~/kernel_module$ lsmod | grep hello
hello                  16384  0
[ 402.783661] Hello World!
[ 568.469739] Bye-bye World!.
```

Task 1.B: Implement a Simple Firewall Using Netfilter

1、和前面一样，将文件拷贝到 /home/seed/ 下面进行编译。

```
[07/26/21]seed@VM:~/packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M] /home/seed/packet_filter/seedFilter.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC [M] /home/seed/packet_filter/seedFilter.mod.o
  LD [M] /home/seed/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/26/21]seed@VM:~/kernel_module$ dig @8.8.8.8 www.example.com
```

```
; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 46429
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                20930   IN      A      93.184.216.34

;; Query time: 259 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Mon Jul 26 02:58:02 EDT 2021
;; MSG SIZE rcvd: 60
```

加载到内核后，可以看到防火墙生效。

```
[07/26/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/26/21]seed@VM:~/packet_filter$ dig @8.8.8.8 www.example.com
```

```
; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached
```

2. 数据报从进入系统, 进行 IP 校验以后, 首先经过第一个 HOOK 函数 NF_INET_PRE_ROUTING 进行处理, 然后就进入路由代码, 其决定该数据报是需要转发还是发给本机的

```
[07/26/21]seed@VM:~/packet_filter$ sudo dmesg -c
[13751.417489] *** PRE_ROUTING
[13751.417491] 10.80.128.28 --> 10.0.2.15 (UDP)
[13751.418269] *** PRE_ROUTING
[13751.418271] 127.0.0.1 --> 127.0.0.53 (UDP)
[13751.421439] *** PRE_ROUTING
[13751.421441] 10.80.128.28 --> 10.0.2.15 (UDP)
[13751.421536] *** PRE_ROUTING
[13751.421537] 127.0.0.53 --> 127.0.0.1 (UDP)
[13784.320530] *** PRE_ROUTING
[13784.320531] 127.0.0.1 --> 127.0.0.1 (UDP)
[13784.320786] *** Dropping 8.8.8.8 (UDP), port 53
[13789.320678] *** Dropping 8.8.8.8 (UDP), port 53
[13794.321599] *** Dropping 8.8.8.8 (UDP), port 53
```

若该数据报是发给本机的, 则该数据经过 HOOK 函数 NF_INET_LOCAL_IN 处理以后然后传递给上层协议。

```
[14207.723137] The filters are being removed.
[14232.637124] Registering filters.
[14235.174264] *** Dropping 8.8.8.8 (UDP), port 53
[14240.175847] *** Dropping 8.8.8.8 (UDP), port 53
[14245.179881] *** Dropping 8.8.8.8 (UDP), port 53
```

若该数据报应该被转发则它被 NF_INET_FORWARD 处理

```
[11011.229107] Registering filters.
[11038.881281] *** LOCAL_OUT
[11038.881282] 127.0.0.1 --> 127.0.0.1 (UDP)
[11038.881702] *** LOCAL_OUT
[11038.881703] 10.0.2.15 --> 8.8.8.8 (UDP)
[11038.881708] *** Dropping 8.8.8.8 (UDP), port 53
[11043.879873] *** LOCAL_OUT
[11043.879875] 10.0.2.15 --> 8.8.8.8 (UDP)
[11043.879884] *** Dropping 8.8.8.8 (UDP), port 53
[11048.879945] *** LOCAL_OUT
[11048.879947] 10.0.2.15 --> 8.8.8.8 (UDP)
[11048.879956] *** Dropping 8.8.8.8 (UDP), port 53
[11051.520558] *** LOCAL_OUT
[11051.520560] 10.0.2.15 --> 10.80.128.28 (UDP)
[11051.525987] *** LOCAL_OUT
[11051.525988] 127.0.0.1 --> 127.0.0.53 (UDP)
[11051.526083] *** LOCAL_OUT
[11051.526083] 10.0.2.15 --> 10.80.128.28 (UDP)
[11051.529714] *** LOCAL_OUT
[11051.529715] 127.0.0.53 --> 127.0.0.1 (UDP)
```

挂载 NF_INET_LOCAL_OUT 时, 本机产生的数据包将会第一个到达此 HOOK, 数据经过

HOOK 函数 NF_INET_LOCAL_OUT 处理后，进行路由选择处理，然后经过 NF_INET_POST_ROUTING 处理后发送出去。

```
[07/26/21]seed@VM:~/packet_filter$ sudo dmesg -c
[14494.411309] The filters are being removed.
[14725.712695] Registering filters.
[14728.213308] *** POST_ROUTING
[14728.213310] 127.0.0.1 --> 127.0.0.1 (UDP)
[14728.214146] *** Dropping 8.8.8.8 (UDP), port 53
[14733.216328] *** Dropping 8.8.8.8 (UDP), port 53
[14738.229683] *** Dropping 8.8.8.8 (UDP), port 53
[14788.992506] *** POST_ROUTING
[14788.992508] 10.0.2.15 --> 35.224.170.84 (TCP)
[14789.995610] *** POST_ROUTING
[14789.995623] 10.0.2.15 --> 35.224.170.84 (TCP)
[14792.011197] *** POST_ROUTING
[14792.011199] 10.0.2.15 --> 35.224.170.84 (TCP)
[14796.139081] *** POST_ROUTING
[14796.139101] 10.0.2.15 --> 35.224.170.84 (TCP)
[14804.331815] *** POST_ROUTING
[14804.331834] 10.0.2.15 --> 35.224.170.84 (TCP)
```

3、修改后的代码如下：

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>
#include <linux/ip.h>
#include <linux/tcp.h>
#include <linux/udp.h>
#include <linux/icmp.h>
#include <linux/if_ether.h>
#include <linux/inet.h>

static struct nf_hook_ops hook1, hook2, hook3, hook4;
unsigned int blockUDP(void *priv, struct sk_buff *skb,
const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct udphdr *udph;
    u16 port = 53;
    char ip[16] = "8.8.8.8";
    u32 ip_addr;
    if (!skb) return NF_ACCEPT;
    iph = ip_hdr(skb);
    // Convert the IPv4 address from dotted decimal to 32-bit binary
    in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);
    if (iph->protocol == IPPROTO_UDP) {
        udph = udp_hdr(skb);
        if (iph->daddr == ip_addr && ntohs(udph->dest) == port){
            printk(KERN_WARNING "*** Dropping %pI4 (UDP), port %d\n", &(iph-
            >daddr), port);
        }
    }
}
```

```

return NF_DROP;
}
}
return NF_ACCEPT;
}
unsigned int blockTCP(void *priv, struct sk_buff *skb,
const struct nf_hook_state *state)
{
struct iphdr *iph;
struct tcphdr *tcph;
u16 port = 23;
char ip[16] = "10.9.0.1";
u32 ip_addr;
if (!skb) return NF_ACCEPT;
iph = ip_hdr(skb);
// Convert the IPv4 address from dotted decimal to 32-bit binary
in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);
if (iph->protocol == IPPROTO_TCP) {
tcph = tcp_hdr(skb);
if (iph->daddr == ip_addr && ntohs(tcph->dest) == port){
printk(KERN_WARNING "*** Dropping %pI4 (TCP), port %d\n", &(iph-
>daddr), port);
return NF_DROP;
}
}
return NF_ACCEPT;
}
unsigned int blockICMP(void *priv, struct sk_buff *skb,
const struct nf_hook_state *state)
{
struct iphdr *iph;
struct icmphdr *icmph;
char ip[16] = "10.9.0.1";
u32 ip_addr;
if (!skb) return NF_ACCEPT;
iph = ip_hdr(skb);
// Convert the IPv4 address from dotted decimal to 32-bit binary
in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);
if (iph->protocol == IPPROTO_ICMP) {
icmph = icmp_hdr(skb);
if (iph->daddr == ip_addr){
printk(KERN_WARNING "*** Dropping %pI4 (ICMP)\n", &(iph->daddr));
return NF_DROP;
}
}
}

```

```

return NF_ACCEPT;
}
unsigned int printInfo(void *priv, struct sk_buff *skb,
const struct nf_hook_state *state)
{
struct iphdr *iph;
char *hook;
char *protocol;
switch (state->hook){
case NF_INET_LOCAL_IN: hook = "LOCAL_IN"; break;
case NF_INET_LOCAL_OUT: hook = "LOCAL_OUT"; break;
case NF_INET_PRE_ROUTING: hook = "PRE_ROUTING"; break;
case NF_INET_POST_ROUTING: hook = "POST_ROUTING"; break;
case NF_INET_FORWARD: hook = "FORWARD"; break;
default: hook = "IMPOSSIBLE"; break;
}
printk(KERN_INFO "*** %s\n", hook); // Print out the hook info
iph = ip_hdr(skb);
switch (iph->protocol){
case IPPROTO_UDP: protocol = "UDP"; break;
case IPPROTO_TCP: protocol = "TCP"; break;
case IPPROTO_ICMP: protocol = "ICMP"; break;
default: protocol = "OTHER"; break;
}
// Print out the IP addresses and protocol
printk(KERN_INFO " %pI4 --> %pI4 (%s)\n",
&(iph->saddr), &(iph->daddr), protocol);
return NF_ACCEPT;
}

int registerFilter(void) {
printk(KERN_INFO "Registering filters.\n");
hook1.hook = printInfo;
hook1.hooknum = NF_INET_LOCAL_OUT;
hook1.pf = PF_INET;
hook1.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook1);
hook2.hook = blockUDP;
hook2.hooknum = NF_INET_POST_ROUTING;
hook2.pf = PF_INET;
hook2.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook2);
hook3.hook = blockICMP;
hook3.hooknum = NF_INET_PRE_ROUTING;
hook3.pf = PF_INET;
hook3.priority = NF_IP_PRI_FIRST;

```



```

nf_register_net_hook(&init_net, &hook3);
hook4.hook = blockTCP;
hook4.hooknum = NF_INET_PRE_ROUTING;
hook4.pf = PF_INET;
hook4.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook4);
return 0;
}

void removeFilter(void) {
    printk(KERN_INFO "The filters are being removed.\n");
    nf_unregister_net_hook(&init_net, &hook1);
    nf_unregister_net_hook(&init_net, &hook2);
    nf_unregister_net_hook(&init_net, &hook3);
    nf_unregister_net_hook(&init_net, &hook4);
}

module_init(registerFilter);
module_exit(removeFilter);
MODULE_LICENSE("GPL");

```

加载内核

```

[07/26/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
seedFilter                16384  0

```

开启容器，在 10.9.0.5 容器上分别进行 ping 10.9.0.1 和 telnet 10.9.0.1 发现都不通过 dmesg 查看：

```

root@89c2a49a4fe2:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
6 packets transmitted, 0 received, 100% packet loss, time 5101ms

root@89c2a49a4fe2:/# telnet 10.9.0.1
Trying 10.9.0.1...
^C
[ 2094.390684] *** Dropping 10.9.0.1 (ICMP), port
[ 2095.396858] *** Dropping 10.9.0.1 (ICMP), port
[ 2096.419797] *** Dropping 10.9.0.1 (ICMP), port
[ 2097.443347] *** Dropping 10.9.0.1 (ICMP), port
[ 2098.466919] *** Dropping 10.9.0.1 (ICMP), port
[ 2099.491227] *** Dropping 10.9.0.1 (ICMP), port
[ 2102.334535] *** Dropping 10.9.0.1 (TCP), port 23
[ 2103.362271] *** Dropping 10.9.0.1 (TCP), port 23
[ 2105.379589] *** Dropping 10.9.0.1 (TCP), port 23

```

Task 2: Experimenting with Stateless Firewall Rules

Task 2.A: Protecting the Router

输入以下命令后， ping 10.9.0.11 和 telnet 10.9.0.11 都不通。

```

root@be47739dff58:/# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@be47739dff58:/# iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCE
PT
root@be47739dff58:/# iptables -P OUTPUT DROP
root@be47739dff58:/# iptables -P INPUT DROP
root@4362d2195d99:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
^C
--- 10.9.0.11 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3075ms

root@4362d2195d99:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C

```

修改成这样，可以 ping 通，但是 telnet 不通。

```

iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEP

iptables -P OUTPUT DROP
iptables -P INPUT DROP

```

设置了 iptables -P OUTPUT DROP 后，二者无法 ping 通，表示丢弃所有外出的包 在单独设置了 iptables -P INPUT DROP，可以发现，router 可以 ping 通其他主机，但是其他主机不可以 通 router，表示所有进入的包都被丢弃了，但是外出的包不受限制

Task 2.B: Protecting the Internal Network

```

iptables -A FORWARD -p icmp --icmp-type echo-request -d 10.9.0.5/24 -j ACCEPT
iptables -A FORWARD -p icmp --icmp-type echo-reply -d 192.168.60.0/24 -j ACCEPT
iptables -A FORWARD -p icmp --icmp-type echo-request -d 192.168.60.0/24 -j ACCEPT
iptables -A INPUT -p icmp -j ACCEPT
iptables -A OUTPUT -p icmp -j ACCEPT
iptables -P FORWARD DROP
Chain INPUT (policy ACCEPT)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         anywhere

Chain FORWARD (policy DROP)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         10.9.0.0/24          icmp echo-request
ACCEPT    icmp -- anywhere         192.168.60.0/24      icmp echo-reply
ACCEPT    icmp -- anywhere         192.168.60.0/24      icmp echo-request

Chain OUTPUT (policy ACCEPT)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         anywhere

```

从外部主机 ping 路由器，可以 ping 通； ping 内部主机不通； telnet 内部主机不通。 内部主机 ping 外部主机，可以 ping 通； telnet 外部主机不通。

Task 2.C: Protecting Internal Servers

```

iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 -j ACCEPT
iptables -A FORWARD -p tcp --sport 23 -s 192.168.60.5 -j ACCEPT
iptables -A FORWARD -d 10.9.0.0/24 -j DROP
iptables -A FORWARD -d 192.168.60.0/24 -j DROP
Chain INPUT (policy ACCEPT)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         anywhere

Chain FORWARD (policy DROP)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         10.9.0.0/24          icmp echo-request
ACCEPT    icmp -- anywhere         192.168.60.0/24      icmp echo-reply
ACCEPT    icmp -- anywhere         192.168.60.0/24      icmp echo-request
ACCEPT    tcp  -- anywhere             host1-192.168.60.5.net-192.168.60.0    tcp dpt:telnet
ACCEPT    tcp  -- host1-192.168.60.5.net-192.168.60.0 anywhere          tcp spt:telnet
DROP      all  -- anywhere             10.9.0.0/24
DROP      all  -- anywhere             192.168.60.0/24

Chain OUTPUT (policy ACCEPT)
target    prot opt source                destination
ACCEPT    icmp -- anywhere         anywhere

```

从外部主机(10.9.0.5)telnet 192.168.60.5 ， 可以连接成功。 从外部主机(10.9.0.5)telnet 192.168.60.6 ， 无法连接。 外部主机不能访问内部服务器， 内部主机可以访问所有内部服务器，

内部主机不可以访问外部服务器 所有内部主机都运行 telnet 服务器(侦听端口 23)。外部主机只能访问 192.168.60.5 上的 telnet 服务器，不能访问其他内部主机。

Task 3: Connection Tracking and Stateful Firewall

Task 3.A: Experiment with the Connection Tracking

ICMP 的连接状态保持时间只有 30 秒左右。

```
root@85267f447274:~# conntrack -L
icmp      1 8 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=31 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=31 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
icmp      1 6 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=31 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=31 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
icmp      1 4 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=31 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=31 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
icmp      1 1 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=31 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=31 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
```

UDP 的连接状态保持时间和也只有 20~30 秒之间。

```
root@85267f447274:~# conntrack -L
tcp        6 53 TIME_WAIT src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
tcp        6 52 TIME_WAIT src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
tcp        6 48 TIME_WAIT src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
tcp        6 44 TIME_WAIT src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
tcp        6 38 TIME_WAIT src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

TCP 的连接状态保持时间非常长，大约 430000 秒。

```
root@85267f447274:~# conntrack -L
tcp        6 431997 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@85267f447274:~# conntrack -L
tcp        6 431994 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=51068 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=51068 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

Task 3.B: Setting Up a Stateful Firewall

在路由器上利用 iptables 命令和连接跟踪机制，创建过滤规则如下：

```
iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 --syn -m conntrack --ctstate NEW -j ACCEPT
iptables -A FORWARD -p tcp --dport 23 -d 10.9.0.0/24 --syn -m conntrack --ctstate NEW -j ACCEPT
iptables -P FORWARD DROP
```

从外部主机(10.9.0.5)telnet 192.168.60.5 可以连接成功。


```
root@597b5efa5bf3:/# telnet 192.168.60.5
```

```
Trying 192.168.60.5...
```

```
Connected to 192.168.60.5.
```

```
Escape character is '^['.
```

```
Ubuntu 20.04.1 LTS
```

```
4644a85d9f8c login: seed
```

```
Password:
```

```
Login incorrect
```

```
4644a85d9f8c login: seed
```

```
Password:
```

```
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

```
* Documentation: https://help.ubuntu.com
```

```
* Management: https://landscape.canonical.com
```

```
* Support: https://ubuntu.com/advantage
```

从外部主机(10.9.0.5)telnet 192.168.0.6 不成功

```
root@597b5efa5bf3:/# telnet 192.168.60.6
```

```
Trying 192.168.60.6...
```

```
^C
```

从内部主机(192.168.60.5)telnet 10.9.0.5 和 192.168.60.6 ， 连接成功。

```
root@4644a85d9f8c:/# telnet 10.9.0.5
```

```
Trying 10.9.0.5...
```

```
Connected to 10.9.0.5.
```

```
Escape character is '^['.
```

```
Ubuntu 20.04.1 LTS
```

```
597b5efa5bf3 login: seed
```

```
Password:
```

```
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

```
* Documentation: https://help.ubuntu.com
```

```
* Management: https://landscape.canonical.com
```

```
Trying 192.168.60.6...
```

```
Connected to 192.168.60.6.
```

```
Escape character is '^['.
```

```
Ubuntu 20.04.1 LTS
```

```
15a777504b7d login: seed
```

```
Password:
```

```
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic
```

```
* Documentation: https://help.ubuntu.com
```

```
* Management: https://landscape.canonical.com
```

Task 4: Limiting Network Traffic

```

root@137e7980c0ce:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.166 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.100 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.104 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.104 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.106 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.101 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.091 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.140 ms
^C
--- 192.168.60.5 ping statistics ---
24 packets transmitted, 8 received, 66.6667% packet loss, time 235
61ms
rtt min/avg/max/mdev = 0.091/0.114/0.166/0.023 ms

```

一开始会以正常速度发送，后面每隔 6s 发一个包。如果只执行第一条命令，会议正常速度发送。

Task 5: Load Balancing

```

iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 0 -j DNAT --to-destination
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 1 -j DNAT --to-destination
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 2 -j DNAT --to-destination

```

按顺序 hello_1 被发送到 192.168.60.5 8080，hello_2 被发送到 192.168.60.6 8080，hello_3 被发送到 192.168.60.7 8080。

```

root@137e7980c0ce:/# echo hello|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello1|nc -u 10.9.0.11 8080
root@137e7980c0ce:/# echo hello_1|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_2|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_3|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_4|nc -u 10.9.0.11 8080
root@137e7980c0ce:/# echo hello_4|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_5|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_6|nc -u 10.9.0.11 8080
root@137e7980c0ce:/# echo hello_6|nc -u 10.9.0.11 8080
^[[A^C
root@137e7980c0ce:/# echo hello_7|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_8|nc -u 10.9.0.11 8080
^C
root@137e7980c0ce:/# echo hello_9|nc -u 10.9.0.11 8080

```

```
root@552e72b0412e:/# hello
hello_1
root@49e64dea0623:/# nc -luk 8080
hello_3
```

```
root@e94e2533f8f6:/# iptables -F
root@e94e2533f8f6:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.33 -j DNAT --to-destination
192.168.60.5:8080
root@e94e2533f8f6:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.33 -j DNAT --to-destination
192.168.60.6:8080
root@e94e2533f8f6:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.33 -j DNAT --to-destination
192.168.60.7:8080
```

虽然是等概率发送数据，但每个主机收到的数量各不相同，甚至有的差异较大，当样本数量足够多时，应该是趋于平均的

```
root@d998244af73c:/# nc -luk 8080
hello_1
hello_4
hello_5
hello_7
hello_9
root@552e72b0412e:/# nc -luk 8080
hello_2
hello_3
hello_8
```

```
root@49e64dea0623:/# nc -luk 8080
hello_6
```